# ML SUPERVISED LEARNING

Ethan Lam

# EDA

- First check of the data

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|---|---|---|---|---|---|---|---|---|---|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 | 31.992578 | 0.471876 | 33.240885 | 0.348958 |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 | 7.884160 | 0.331329 | 11.760232 | 0.476951 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.078000 | 21.000000 | 0.000000 |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 | 27.300000 | 0.243750 | 24.000000 | 0.000000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 | 32.000000 | 0.372500 | 29.000000 | 0.000000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 | 36.600000 | 0.626250 | 41.000000 | 1.000000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 | 67.100000 | 2.420000 | 81.000000 | 1.000000 |

# EDA – HANDLING MISSING VALUES

- Update missing values by median
- No duplicated values found
- Outliers has been founded in Insulin feature

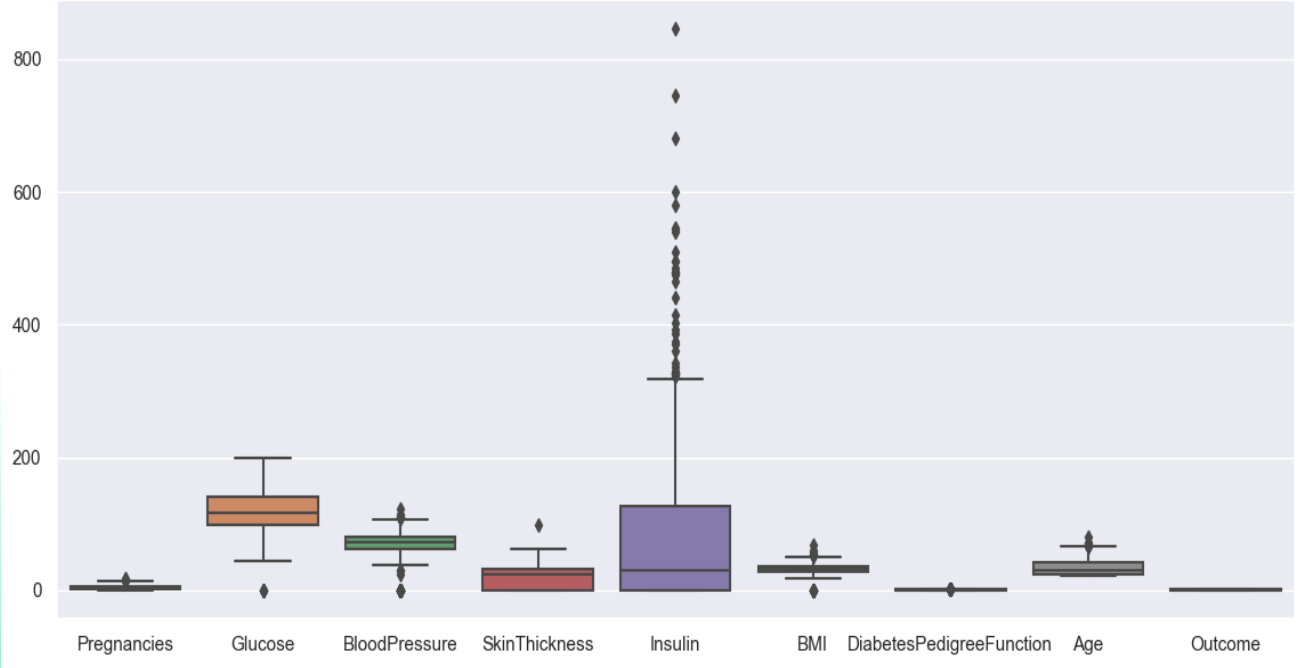

# SCALING

- Scaling is necessary because the features in dataset have different scales
- The output of the scaling

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|---|---|---|---|---|---|---|---|
| 0 | 0.639947 | 0.848324 | 0.149641 | 0.907270 | -0.692891 | 0.204013 | |
| 1 | -0.844885 | -1.123396 | -0.160546 | 0.530902 | -0.692891 | -0.684422 | |
| 2 | 1.233880 | 1.943724 | -0.263941 | -1.288212 | -0.692891 | -1.103255 | |
| 3 | -0.844885 | -0.998208 | -0.160546 | 0.154533 | 0.123302 | -0.494043 | |
| 4 | -1.141852 | 0.504055 | -1.504687 | 0.907270 | 0.765836 | 1.409746 | |
| .. | ... | ... | ... | ... | ... | ... | |
| 763 | 1.827813 | -0.622642 | 0.356432 | 1.722735 | 0.870031 | 0.115169 | |
| 764 | -0.547919 | 0.034598 | 0.046245 | 0.405445 | -0.692891 | 0.610154 | |
| 765 | 0.342981 | 0.003301 | 0.149641 | 0.154533 | 0.279594 | -0.735190 | |
| 766 | -0.844885 | 0.159787 | -0.470732 | -1.288212 | -0.692891 | -0.240205 | |
| 767 | -0.844885 | -0.873019 | 0.046245 | 0.656358 | -0.692891 | -0.202129 | |

| | DiabetesPedigreeFunction | Age |
|---|---|---|
| 0 | 0.468492 | 1.425995 |
| 1 | -0.365061 | -0.190672 |
| 2 | 0.604397 | -0.105584 |
| 3 | -0.920763 | -1.041549 |
| 4 | 5.484909 | -0.020496 |
| .. | ... | ... |
| 763 | -0.908682 | 2.532136 |
| 764 | -0.398282 | -0.531023 |
| 765 | -0.685193 | -0.275760 |
| 766 | -0.371101 | 1.170732 |
| 767 | -0.473785 | -0.871374 |

# NORMALIZATION

```python
# # Normalization
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df_train_normalized = scaler.fit_transform(df[num_feats])

print(df_train_normalized)
```

✓ 0.0s

```
[[0.375      0.67096774 0.5        ... 0.41621622 0.41895604 0.56862745]
 [0.0625     0.26451613 0.425      ... 0.22702703 0.20833333 0.19607843]
 [0.5        0.89677419 0.4        ... 0.13783784 0.4532967  0.21568627]
 ...
 [0.3125     0.49677419 0.5        ... 0.21621622 0.127442   0.17647059]
 [0.0625     0.52903226 0.35       ... 0.32162162 0.20680708 0.50980392]
 [0.0625     0.31612903 0.475      ... 0.32972973 0.18086081 0.03921569]]
```

# TRAINING MODEL

- There are three models used

  Linear regression

  XGBRegressor

  RandomForestRegressor

- The model evaluation metrics

  Mean squared error

  F1 score

  AUC-score

→ Model Random Forest has the best fit as it has the higher F1Score and AUC score

# CONCLUSION

- Linear regression is the best fit model in this scenario
- The most difficult part is to do feature selection